

# X-RAY Bone Classification

With MURA database and DenseNet169 pre-trained on imagenet  
By: Aileen Dugan

# Setup of Data with pandas dataframes

Extract only the shoulder x-ray images for both training images and then testing images using the .csv files provided by MURA

```
[ ] 1 #A function to generate the dataframe for a csv file
    2 def generate_df(dataset_root, csv_name, BODYPART):
    3     df = pd.read_csv(dataset_root/csv_name, header=None, names=['filename'])
    4     df['class'] = (df.filename
    5                   .str.extract('study.*_(positive|negative)'))
    6     df['BodyPart'] = (df.filename
    7                      .str.extract('XR_(SHOULDER|ELBOW|FINGER|FOREARM|HAND|HUMERUS|WRIST)'))
    8     bodypart = df[(df["BodyPart"]==BODYPART)]
    9     df_onlyOne = bodypart[["filename","class"]]
   10     return df_onlyOne
```

```
▶ 1 df_train = generate_df(dataset_root, 'train_image_paths.csv',"SHOULDER")
   2 df_train.head()
```



	filename	class
0	MURA-v1.1/train/XR_SHOULDER/patient00001/study...	positive
1	MURA-v1.1/train/XR_SHOULDER/patient00001/study...	positive
2	MURA-v1.1/train/XR_SHOULDER/patient00001/study...	positive
3	MURA-v1.1/train/XR_SHOULDER/patient00002/study...	positive
4	MURA-v1.1/train/XR_SHOULDER/patient00002/study...	positive



# Turn DataFrames into ndarrays of images and labels

Get images from dataframe filepaths and size them properly to stack into an ndarray. Also create an array of labels using the dataframe's class of positive or negative for each image


```
1 img_path = '/content/drive/MyDrive/MURA/'

1 #take dataframe and make ndarray items for training and testing - fit to densenet keep 3 channels wit
2 x_train = []
3 y_train = []
4 for index,row in df_train.iterrows():
5     input_arr = cv2.imread(os.path.join(img_path,row["filename"]))
6     #, cv2.IMREAD_GRAYSCALE)
7     input_arr = cv2.resize(input_arr, (224,224))
8     label = row["class"]
9     x_train.append(input_arr)
10    y_train.append(label)
11 x_train = np.stack(x_train, axis=0)
12 y_train = np.array(y_train)
```

Show an example image to make sure this was done properly - shows image and it's label

```
1 # Let's show one example from the dataset
2 image_index = 688 # you may select anything from 0 to 8378 (we have 8379 training samples for shoulder)
3 print(y_train[image_index])
4 plt.imshow(x_train[image_index], cmap='Greys')
5 plt.show()
```

positive



# Reshape ndarrays to fit for imagenet as 4-dimensional with 3 layers, float32 type, and 8-bit grayscale intensities

```
▶ 1 # Reshaping the array to 4-dims so that it can work with the Keras API
2 x_train = x_train.reshape(x_train.shape[0], 224, 224, 3)
3 x_test = x_test.reshape(x_test.shape[0], 224, 224, 3)
4 input_shape = (224, 224, 3)
5
6 # Making sure that the values are float so that we can get decimal points after division
7 x_train = x_train.astype('float32')
8 x_test = x_test.astype('float32')
9
10 # Normalizing the 8-bit grayscale intensities by dividing them by the max intensity value
11 x_train /= 255
12 x_test /= 255
```

# Download and build the DenseNet169 Architecture from tf.keras

```
[ ] 1 #Downloading the densenet model pretrained on the imagenet dataset
    2 densenet = tf.keras.applications.DenseNet169(weights='imagenet', include_top = False, input_shape=(224, 224, 3))


Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/densenet/densenet169\_weights\_tf\_dim\_ordering\_tf\_data\_format.h5
51877672/51877672 [=====] - 0s 0us/step
```

```
[ ] 1 #Freezing the weights of the pretrained model
    2 densenet.trainable = False
```

```
[ ] 1 densenet.summary()
```

conv5_block20_1_relu (Activation)	(None, 7, 7, 128)	0	['conv5_block20_1_bn[0][0]']
conv5_block20_2_conv (Conv2D)	(None, 7, 7, 32)	36864	['conv5_block20_1_relu[0][0]']
conv5_block20_concat (Concatenate)	(None, 7, 7, 1280)	0	['conv5_block19_concat[0][0]', 'conv5_block20_2_conv[0][0]']
conv5_block21_0_bn (Batch Normalization)	(None, 7, 7, 1280)	5120	['conv5_block20_concat[0][0]']
conv5_block21_0_relu (Activation)	(None, 7, 7, 1280)	0	['conv5_block21_0_bn[0][0]']

# DenseNet169 Architecture specs



Layers	Output Size	DenseNet-121	DenseNet-169	DenseNet-201	DenseNet-264
Convolution	$112 \times 112$	$7 \times 7$ conv, stride 2			
Pooling	$56 \times 56$	$3 \times 3$ max pool, stride 2			
Dense Block (1)	$56 \times 56$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$
Transition Layer (1)	$56 \times 56$	$1 \times 1$ conv			
	$28 \times 28$	$2 \times 2$ average pool, stride 2			
Dense Block (2)	$28 \times 28$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$
Transition Layer (2)	$28 \times 28$	$1 \times 1$ conv			
	$14 \times 14$	$2 \times 2$ average pool, stride 2			
Dense Block (3)	$14 \times 14$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 24$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 48$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 64$
Transition Layer (3)	$14 \times 14$	$1 \times 1$ conv			
	$7 \times 7$	$2 \times 2$ average pool, stride 2			
Dense Block (4)	$7 \times 7$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 16$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 48$
Classification Layer	$1 \times 1$	$7 \times 7$ global average pool			
		1000D fully-connected, softmax			

Table 4: DenseNet Architecture

# Binary Cross entropy loss and Adam optimizer for DenseNet169

```
1 #Compiling the model using adam optimizer
2 model.compile(loss=tf.keras.losses.binary_crossentropy,
3               optimizer=tf.keras.optimizers.Adam(),
4               metrics=['accuracy'])
```



# Encode Categorical labels (Positive/Negative) as numerical values 1 and 0

```
[ ] 1 from sklearn import preprocessing
    2 # prepare labels so that positive and negative map to 0 and 1
    3 def prepare_targets(y_train, y_test):
    4     le = preprocessing.LabelEncoder()
    5     le.fit(y_train)
    6     y_train_enc = le.transform(y_train)
    7     y_test_enc = le.transform(y_test)
    8     return y_train_enc, y_test_enc
```

```
[ ] 1 def decode_targets(y_train, prep_label):
    2     le = preprocessing.LabelEncoder()
    3     le.fit(y_train)
    4     label = le.inverse_transform(prep_label)
    5     #y_test_enc = le.transform(y_test)
    6     return label
```

```
[ ] 1 y_train_enc, y_test_enc = prepare_targets(y_train, y_test)
```

# Train denseNet169 on 8379 training images with 5 epochs

- accuracy of ~83%

```
+ Code + Text
[ ] 1 epochs = 5
    2 print('-TRAINING-----')
    3 print('Input shape:', x_train.shape)
    4 print('Number of training images: ', x_train.shape[0])
    5
    6 model.fit(x=x_train, y=y_train_enc, epochs=epochs)

--TRAINING-----
Input shape: (8379, 224, 224, 3)
Number of training images: 8379
Epoch 1/5
262/262 [=====] - 1574s 6s/step - loss: 1.1167 - accuracy: 0.6484
Epoch 2/5
262/262 [=====] - 1555s 6s/step - loss: 0.8418 - accuracy: 0.7334
Epoch 3/5
262/262 [=====] - 1550s 6s/step - loss: 0.6444 - accuracy: 0.7894
Epoch 4/5
262/262 [=====] - 1502s 6s/step - loss: 0.4345 - accuracy: 0.8417
Epoch 5/5
262/262 [=====] - 1518s 6s/step - loss: 0.4868 - accuracy: 0.8385
<keras.callbacks.History at 0x7f38226362e0>
```

# Evaluate on testing images accuracy ~67%

```
1 print('-TESTING-----')
2 print('Number of test images:', x_test.shape[0])
3 score = model.evaluate(x_test, y_test_enc)
4 print('Test loss:', score[0])
5 print('Test accuracy:', score[1])
6
7 # Print 10 example test digits with their true and predicted labels
8 fig, axes = plt.subplots(2, 5)
9 fig.tight_layout(rect=(0,0,3,3))
10
11 image_idx = np.random.randint(1,562,(2,5))
12
13 for i, j in it.product(range(2), range(5)):
14     test_image = x_test[image_idx[i,j]].reshape(1, 224, 224, 3)
15     test_label = y_test[image_idx[i,j]]
16     sigmoid_outputs = model.predict(test_image)
17     pred_label = sigmoid_outputs.argmax()
18     pred_label = decode_targets(y_train,[pred_label])
19
20     axes[i, j].imshow(test_image.reshape(224, 224,3),cmap='Greys')
21     axes[i, j].set_aspect('equal', 'box')
22     axes[i, j].set_title("actual: {} predicted: {}".format(test_label,pred_label[0]))
23
24 plt.show()
```

# Evaluate on testing images accuracy ~67% 10 example images

```
-TESTING-----  
Number of test images: 563  
18/18 [=====] - 107s 6s/step - loss: 1.3024 - accuracy: 0.6785  
Test loss: 1.3024436235427856  
Test accuracy: 0.6785079836845398  
1/1 [=====] - 3s 3s/step  
1/1 [=====] - 0s 289ms/step  
1/1 [=====] - 0s 272ms/step  
1/1 [=====] - 0s 280ms/step  
1/1 [=====] - 0s 279ms/step  
1/1 [=====] - 0s 265ms/step  
1/1 [=====] - 0s 268ms/step  
1/1 [=====] - 0s 259ms/step  
1/1 [=====] - 0s 259ms/step  
1/1 [=====] - 0s 257ms/step
```

